

Time Series Notes

Nathan Ueda

April 24, 2024

Contents

1	Time Series Basics	4
2	Stationarity and Ergodicity	4
2.1	Basics and Importance	4
2.2	Strongly Stationarity and Weakly Stationarity	4
2.3	Testing for Stationarity	5
2.4	Ergodicity	6
2.5	Law of Large Numbers and Central Limit Theorem for an Autocorrelated Process	6
3	White Noise	7
4	Autocovariance and Autocorrelation	7
4.1	Autocovariance	7
4.2	Autocorrelation	8
4.3	Testing for autocorrelation	8
5	Lag and Difference Operators	9
5.1	Lag Operators	9
5.2	Difference Operators	10
6	Autoregressive Processes	10
6.1	AR(1) Models	10
6.2	AR(p) Models	11
7	Moving Average Processes	12
7.1	MA(1)	12
7.2	MA(q)	12
8	ARMA Processes	12
8.1	ARMA(p, q)	12
9	ARIMA Processes	12
10	Impulse Response Functions	13
11	Unit Root Tests	13
11.1	Dickey-Fuller Test	14
11.2	Augmented Dickey-Fuller Test	14
12	Regression	15
13	Estimation for AR Models	15
13.1	OLS for AR(p) Process	15
13.2	Autocorrelated Residuals	16
14	Spurious Regression	16

15 Estimation	16
15.1 Estimation Basics	16
15.2 Properties of Estimation	16
15.3 Ordinary Least Squares	16
15.4 Generalized Least Squares	17
15.5 Heteroskedasticity and Autocorrelation Consistent (HAC) Standard Errors	17
15.6 Maximum Likelihood	18
15.7 MLE with Autocorrelation	18
16 Multivariate Time Series	18
16.1 Vector Autoregressions (VAR)	19
16.1.1 VAR(1)	19
16.1.2 VAR(p)	19
16.1.3 Cross-Correlation Function	20
16.1.4 Forecasting	20
16.2 Granger Causality	20
17 Cointegration	20
17.1 Cointegration with Two Time Series	20
18 Vector Error Correction Models	21
19 Present Value Relationships	21
19.1 Gordon Growth Model	21
20 CAPM	22
21 Misc.	22
22 Definitions	22

1 Time Series Basics

What is a time series?

- A *time series* is a sequence of observations of the same variable indexed in time order (i.e. monthly stock returns).
- Defining x_t as a r.v., a time series may be written

$$\{x_1, x_2, \dots, x_T\} \text{ or } \{x_t\}, t = 1, 2, \dots, T$$

Modeling time series data

- The models impose structure, so when dealing with model selection, it is important to evaluate to see if the model captures the features you believe to be present in the data.

2 Stationarity and Ergodicity

2.1 Basics and Importance

- Time series are typically not i.i.d. (i.e. if GNP is unusually high today, GNP will likely be unusually high tomorrow). So instead, we need different desirable properties for time series data. Two of the most important of these properties are stationarity and ergodicity.
- Stationarity is a property for time series with some time-invariant behavior.
- Ergodicity is a property for times series that expresses the idea that the effect of the past on the future eventually dies out.
- Time series with these properties are easier to estimate.
- Time series that are nonstationarity and nonergodic require a different set of techniques.
- Under the assumption of i.i.d r.v.s, we had LLN and CLT. For r.v.s that are not i.i.d. and are instead autocorrelated, stationarity and ergodicity are similar results.
- In practice, by saying the time series is stationary, it is typically implied it is both stationary and ergodic.
- If a times series is stationary, it does not imply ergodicity; however, a time series that is ergodic is always stationary.
- Takeaway: Ensure your time series is stationary and ergodic.

2.2 Strongly Stationarity and Weakly Stationarity

- A process $\{x_t\}$ is *strongly stationary* or *strictly stationary* if all aspects of its behavior are unchanged by shifts in time. More formally, it is defined as the requirement that for every m and n , the distributions of $\{x_1, \dots, x_n\}$ and $\{x_{1+m}, \dots, x_{n+m}\}$ are the same, that is, the joint probability distribution of a sequence of n observations does not depend on their time origin.

- A process $\{x_t\}$ is *weakly stationary* or *covariance stationary* if its mean, variance, and covariance are unchanged by time shifts. More formally, it is defined as weakly stationary if
 - First moment is a finite constant: $E(x_t) = \mu$
 - Second moment is a finite constant: $Var(x_t) = \sigma^2$
 - $Cov(x_t, x_s) = \gamma(|t - s|)$

In other words, the mean and variance do not change with time and the covariance between two observations depends on the time distance between them (aka having same number of observations), not the specific points.

- Strong stationarity does not imply weak stationarity and weak stationarity does not imply strong stationarity.
- Stationarity is important as a stationary process can be modeled with relatively few parameters.

2.3 Testing for Stationarity

- When a times series is observed, a natural question is whether it appears to be stationary.
- **Time series plot:** Looking at a *time series plot* (plot of the series in chronological order) may be useful. If the time series is a stationary series, it should show some random oscillation around some fixed level, a phenomenon called *mean reversion*. If the series wanders without returning repeatedly to some fixed level, then the series should not be modeled as a stationary process.




Figure 1: Stationary vs. non-stationary time series

- If all the roots of a characteristic polynomial for an AR process are greater than 1 in absolute value, then the AR process is stationary.
- Any finite-order MA(q) process is stationary and ergodic.

2.4 Ergodicity

- A stochastic process $\{x_t\}$ is ergodic if any two random variables sufficiently far apart in time are essentially independent. In other words, $\{x_t\}$ is ergodic if x_t and x_{t-j} are close to uncorrelated if j is large enough.

2.5 Law of Large Numbers and Central Limit Theorem for an Autocorrelated Process

- The *ergodic theorem* is a LLM for an autocorrelated process that states, if $\{x_t\}$ is stationary and ergodic, then

$$\bar{x}_T = \frac{1}{T} \sum_{t=1}^T x_t \rightarrow E(x_t) \text{ as } T \rightarrow \infty$$

- If $\{x_t\}$ is stationary and ergodic, then the asymptotic distribution of the sample mean is normal.

3 White Noise




Figure 2: White noise

- The building block for time series models is the *white noise process*, denoted ϵ_t .
- In the least general case,

$$\epsilon_t \sim \text{i.i.d. } N(0, \sigma_\epsilon^2)$$

The assumption of i.i.d. have the following implications

1. No predictability: Past values of a white noise process contain no information to predict future values. Therefore, the best predictor is its mean, which is the same prediction without observing past values. More formally,

$$E(\epsilon_t) = E(\epsilon_t | \epsilon_{t-1}, \epsilon_{t-2} \dots) = 0$$

2. No autocorrelation: Each observation is independent each other. More formally,

$$E(\epsilon_t \epsilon_{t-j}) = \text{Cov}(\epsilon_t \epsilon_{t-j}) = 0$$

3. Conditional homoskedasticity: The conditional variance is a constant, and is the same variance without observing past values. More formally,

$$\text{Var}(\epsilon_t) = \text{Var}(\epsilon_t | \epsilon_{t-1}, \epsilon_{t-2} \dots) = \sigma_\epsilon^2$$

4 Autocovariance and Autocorrelation

4.1 Autocovariance

- Autocovariance: Specifies the covariance between the value of a process at two times.

- Covariance: A nonstandardized measure to quantify the relationship between two variables. Can take on any value. A positive value means the variables tend to move in the same direction, a negative values means the variables tend to move in the opposite direction, and a zero value means they are independent and don't move in relation to each other.
- The *autocovariance* of a series x_t is defined as

$$\gamma_j = \text{Cov}(x_t, x_{t-j})$$

- $\gamma(h) = \gamma(-h)$ since what is important is the the space between the two observations, rather than the exact observations themselves.
- $\gamma_0 = \sigma^2$

4.2 Autocorrelation

- Autocorrelation (serial correlation): The degree of correlation of the same variable between different time intervals.
- Correlation: A standardized measure to quantify the relationship between two variables. Can take on a value inclusively between -1 and 1.
- A time series with autocorrelation implies that, predictive power (i.e. knowing the price of a stock today helps forecast its price tomorrow).
- The *autocorrelation* of a series x_t is defined as

$$\rho_j = \frac{\gamma_j}{\gamma_0}$$

4.3 Testing for autocorrelation

ACF Plots

- Show a correlation between a time series and lagged versions of itself.
- The plot also includes test bounds used to test the null hypothesis that an autocorrelation coefficient is 0. The null is rejected if the sample autocorrelation is outside the bounds, The usual level of the test is 0.05, so one can expect to see about 1 out of 20 samples outside the bounds simply by chance.

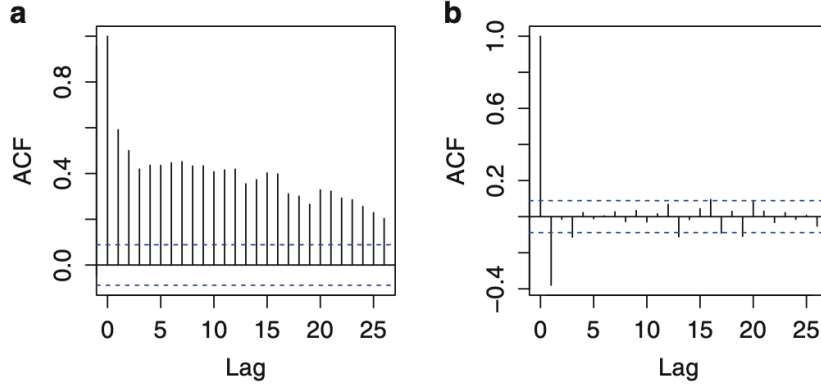


Figure 3: Sample ACF plots of (a) one-month inflation rate (decays slowly, indicating either nonstationarity or long-memory dependence) and (b) changes in the inflation rate (decays to 0 quickly, indicating that it is stationary).

Ljung-Box Test

- A *simultaneous test* is one that tests whether a group of null hypotheses are all true versus the alternative that at least one of them is false.
- The *Ljung-Box test* is a simultaneous test where the null is $\rho_1 = \rho_2 = \dots = \rho_K$ for some K .
- If the Ljung-Box test rejects, then we conclude that one or more of $\rho_1, \rho_2, \dots, \rho_K$ is nonzero.

5 Lag and Difference Operators

- These operators are useful in describing some time series models.
- They allow us to simplify notation.

5.1 Lag Operators

- The lag operator moves the index back one unit in time.
- The lag operator L is defined

$$Lx_t = x_{t-1}$$

$$L^j x_t = x_{t-j}$$

- Also commonly referred to as a backwards operator.
- Given an $ARMA(p, q)$ model, we could simplify it as follows

$$\begin{aligned} Y_t &= \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \phi_p Y_{t-p} + \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots + \theta_q \epsilon_{t-q} \\ Y_t - \phi_1 Y_{t-1} - \phi_2 Y_{t-2} - \dots - \phi_p Y_{t-p} &= \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \dots + \theta_q \epsilon_{t-q} \\ Y_t - \phi_1 L Y_t - \phi_2 L^2 Y_t - \dots - \phi_p L^p Y_t &= \epsilon_t + \theta_1 L \epsilon_t + \theta_2 L^2 \epsilon_t + \dots + \theta_q L^q \epsilon_t \\ (1 - \phi_1 L - \phi_2 L^2 - \dots - \phi_p L^p) Y_t &= (1 + \theta_1 L + \theta_2 L^2 + \dots + \theta_q L^q) \epsilon_t \end{aligned}$$

5.2 Difference Operators

- Defined as $\Delta = 1 - L$ so that

$$\Delta x_t = x_t - Lx_t = x_t - x_{t-1}$$

- Δ^k is called the k -th order differencing operator.

$$\Delta^k x_t = (1 - L)^k x_t$$

6 Autoregressive Processes

- Time series models with correlation can be constructed from white noise.
- The simplest correlated stationary processes are *autoregressive processes*, where $\{x_t\}$ is modeled as a weighted average of past observations plus a white noise “error.”
- The term *autoregression* refers to the regression of the process on its own past values.
- The ACF of AR processes declines geometrically.

6.1 AR(1) Models

- Let $\epsilon_1, \epsilon_2, \dots$ be $\text{WN}(0, \sigma_\epsilon^2)$. Then $\{x_t\}$ is an AR(1) process if, for some constant parameter ϕ

$$x_t = \phi x_{t-1} + \epsilon_t$$

for all t .

- ϕx_{t-1} may be thought of as representing the memory of the past observation into the present value of the process. This is what we believe we can model and predict.
- ϕ determines the amount of feedback, where a larger absolute value results in more feedback.
- ϵ_t represents the effect of new information that cannot be modeled, hence, it is represented as white noise. This is what we cannot model nor predict.
- Properties of a Stationary AR(1) Process

$$- E(x_t) = \mu$$

$$- \text{Var}(x_t) = \gamma_0 = \sigma_x^2 = \frac{\sigma_\epsilon^2}{1 - \phi^2}$$

- The ACF of an AR(1) process depends only on one parameter, ϕ . This parsimony comes at the cost that the ACF has only a very limited range of shapes. If the ACF does not behave in one of these shapes, the AR(1) model is not suitable.

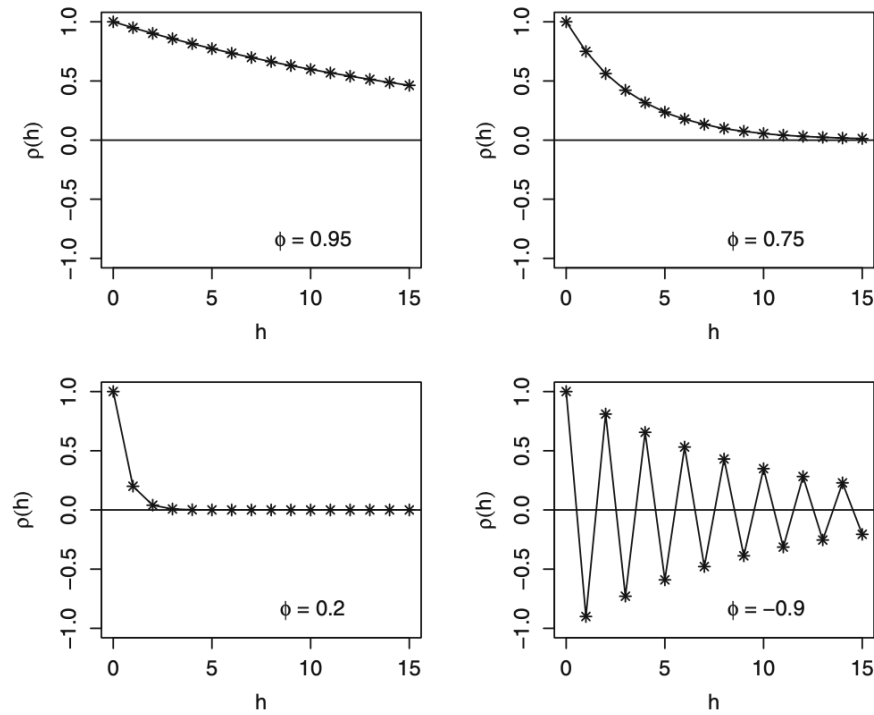


Figure 4: ACF of AR(1) processes with ϕ equal to 0.95, 0.75, 0.2, and -0.9.

- When $|\phi| < 1$ the process is stationary.
- When $\phi = 1$ the process is not stationary and takes the form of a random walk. Recall a random walk, in each period, takes a step random step that is i.i.d. from its previous steps.
- When $|\phi| > 1$ the process is non-stationary and has explosive behavior.
- A non-zero value of ϕ mean that there is some information in the previous observation, but a small value of ϕ means the prediction will not be very accurate.

6.2 AR(p) Models

- A more flexible adaptation of an AR model that is still parsimonious and regresses on the p past values.
- Let $\epsilon_1, \epsilon_2, \dots$ be $WN(0, \sigma_\epsilon^2)$. Then $\{x_t\}$ is an AR(p) process if, for constant parameters $\phi_1, \phi_2, \dots, \phi_{t-p}$

$$x_t = \phi_1 x_{t-1} + \phi_2 x_{t-2} + \dots + \phi_p x_{t-p} + \epsilon_t$$

- If the AR(p) model fits the time series well, then the residuals should look like white noise. The residual autocorrelation can be detected examining the sample ACF of the residuals and using the Ljung-Box test. Any significant residual autocorrelation is a sign the AR(p) model does not fit well.
- A problem with AR models is that they often need a rather large value of p to fit a dataset.

7 Moving Average Processes

- With AR models, feeding past values into the current value has the effect of having at least some correlation at all lags. Sometimes data show correlation only at short lags, in these cases a MA process may be a suitable alternative.
- A process $\{x_t\}$ is a *moving average process* if $\{x_t\}$ can be expressed as a weighted average (moving average) of the past values of the white noise process $\{\epsilon_t\}$.

7.1 MA(1)

- The MA(1) (moving average of order 1) process is

$$x_t = \epsilon_t + \theta\epsilon_{t-1}$$

where, as before, the ϵ_t are weak $\text{WN}(0, \sigma_\epsilon^2)$.

7.2 MA(q)

- The MA(q) (moving average of order q) process is

$$x_t = \epsilon_t + \theta_1\epsilon_{t-1} + \theta_2\epsilon_{t-2} + \dots + \theta_q\epsilon_{t-q}$$

where, as before, the ϵ_t are weak $\text{WN}(0, \sigma_\epsilon^2)$.

8 ARMA Processes

- Stationary time series with complex autocorrelation behavior are often more parsimoniously modeled by mixed AR and MA (ARMA) processes rather than by a pure AR or pure MA process.

8.1 ARMA(p, q)

- An ARMA(p, q) model combines both AR and MA terms, and is defined by the equation

$$x_t = \phi_1x_{t-1} + \dots + \phi_px_{t-p} + \epsilon_t + \theta_1\epsilon_{t-1} + \dots + \theta_q\epsilon_{t-q}$$

which shows how $\{x_t\}$ depends on lagged values of itself and lagged values of the white noise process.

9 ARIMA Processes

- Often the first or perhaps second differences of nonstationarity time series are stationary. *Autoregressive integrated moving average* (ARIMA) processes include stationary as well as nonstationary processes.
- A time series $\{x_t\}$ is said to be an ARIMA(p, d, q) process if $\Delta^d x_t$ is ARMA(p, q).
- An ARIMA(p, d, q) is stationary if $d = 0$, otherwise its difference of order d or above are stationary.

- An $\text{ARIMA}(p, 0, q)$ is the same as an $\text{ARMA}(p, q)$.
- A process is $I(d)$ if it is stationary after being differenced d times.
- An $\text{ARIMA}(p, d, q)$ process has d unit roots, therefore we want to difference the process d times to get rid of the unit roots.

10 Impulse Response Functions

- The idea of an IRF is to enact a single shock to ϵ_t (in period t) and, via the IRF, see how the shock affects x_{t+1}, x_{t+2}, \dots
- It allows us to start thinking about causes and effects.
- If the impulse response is sizable for even long horizons indicates possible nonstationarity.
- For an $\text{AR}(1)$ process, notice that the effect of the shock never dies when $\phi = 1$, and it dies out quicker and quicker as we move from $\phi = 0.95$ to $\phi = 0.5$. This makes sense as the closer to 1 $|\phi|$ is, the more useful information there is in the previous observation in forecasting the current observation.




Figure 5: IRF on an $\text{AR}(1)$ Process

11 Unit Root Tests

- Determining whether a time series is best modeled as stationary or nonstationary can be difficult, unit root tests aid in this process.
- Recall the definition of an $\text{ARMA}(p, q)$ process

$$x_t = \phi_1 x_{t-1} + \dots + \phi_p x_{t-p} + \epsilon_t + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q}$$

- The condition of $\{x_t\}$ to be stationary is that all roots of the polynomial

$$1 - \phi_1 x - \dots - \phi_p x^p$$

have absolute values greater than one.

- If there is a unit root, that is, a root with an absolute value equal to 1, then the ARMA process is nonstationary and behaves much like a random walk. This is called the unit root case.
- Unit root tests are used to decide if an AR model have an absolute root equal to 1.
- We can always look at a time series and say we think it is or isn't stationary, but the Dickey-Fuller and Augmented Dickey-Fuller tests give us a robust way to do so.
- A popular unit root test is the augmented Dickey-Fuller test (ADF test). In this test
 - H_0 : there is a unit root (the process is nonstationary)
 - H_1 : the process is stationary

11.1 Dickey-Fuller Test

- Assumes our time series is an $AR(1)$ model, that is

$$Y_t = \phi Y_{t-1} + \epsilon_t$$

- H_0 : there is a unit root, that is, $\phi = 1$ (the process is nonstationary)
- H_1 : the process is stationary, that is $|\phi| < 1$
- Let $\pi = (\phi - 1)$. The Dickey-Fuller test rewrites the $AR(1)$ model as follows to make the left hand side stationary.

$$\begin{aligned} Y_t &= \phi Y_{t-1} + \epsilon_t \\ Y_t - Y_{t-1} &= \phi Y_{t-1} - Y_{t-1} + \epsilon_t \\ \Delta Y_t &= (\phi - 1)Y_{t-1} + \epsilon_t \\ \Delta Y_t &= \pi Y_{t-1} + \epsilon_t \end{aligned}$$

- We now rewrite the hypotheses in terms of π , that is
 - H_0 : $\pi = 0$
 - H_1 : $\pi < 0$

11.2 Augmented Dickey-Fuller Test

- Assumes our time series is an $AR(p)$ model, that is

$$Y_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \phi_p Y_{t-p} + \epsilon_t$$

- H_0 : there is a unit root, that is, $\phi = 1$ (the process is nonstationary)
- H_1 : the process is stationary, that is $|\phi| < 1$

- Let $\pi = (\phi - 1)$. The Dickey-Fuller test rewrites the $AR(p)$ model as follows to make the left hand side stationary.

$$Y_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \dots + \phi_p Y_{t-p} + \epsilon_t$$

$$\Delta Y_t = \beta_0 + \beta_1 t + \pi Y_{t-1} + \sum_{j=2}^p \gamma_j \Delta Y_{t-j} + \epsilon_t$$

- We now rewrite the hypotheses in terms of π , that is

- $H_0: \pi = 0$
- $H_1: \pi < 0$

12 Regression

13 Estimation for AR Models

- AR models can be estimated by OLS. Recall OLS is defined

$$y = X\beta + e$$

where

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \in \mathbb{R}^{n \times 1}, X = \begin{bmatrix} x'_1 \\ x'_2 \\ \vdots \\ x'_n \end{bmatrix} \in \mathbb{R}^{n \times k}, e = \begin{bmatrix} e_1 \\ e_2 \\ \vdots \\ e_n \end{bmatrix} \in \mathbb{R}^{n \times 1}$$

where k is the number of regressors.

- Since true errors are not observed, MA models (which depend upon the true errors), MA models are not estimated by OLS.

13.1 OLS for $AR(p)$ Process

- Recall the definition of an $AR(p)$ process

$$z_t = \alpha + \phi_1 z_{t-1} + \phi_2 z_{t-2} + \dots + \phi_p z_{t-p} + \epsilon_t, \quad t = 1, \dots, T$$

- We let

$$y_i = z_t$$

$$x_i = [1, z_{t-1}, \dots, z_{t-p}]'$$

$$\beta = [\alpha, \phi_1, \dots, \phi_p]'$$

where

$$y = \begin{bmatrix} z_{p+1} \\ z_{p+2} \\ \vdots \\ z_T \end{bmatrix}, X = \begin{bmatrix} 1 & z_p & \dots & z_1 \\ 1 & z_{p+1} & \dots & z_2 \\ \vdots & \vdots & \ddots & \vdots \\ 1 & z_{T-1} & \dots & z_{T-p} \end{bmatrix}, e = \begin{bmatrix} e_1 \\ e_2 \\ \vdots \\ e_n \end{bmatrix}$$

13.2 Autocorrelated Residuals

- If the residuals are autocorrelated, the regression is misspecified.

14 Spurious Regression

- This is concluding there is a relationship between variables when there is none.
- The regression is spurious when we regress one random walk onto another independent random walk. It is spurious as the regression will most likely indicate a non-existent relationship.
- Rule of thumb: If we have an R^2 greater than our Durbin Watson statistic, our relationship is likely spurious.
- Always check the stationarity of the residual. The regression is spurious if the residual is nonstationary (cannot reject the null of the unit root test).
- Just because two series move together does not mean they are related.

15 Estimation

15.1 Estimation Basics

- An estimator is any procedure or formula that is used to predict or estimate the true value of some unknown quantity by using information contained in data points of a sample.

15.2 Properties of Estimation

- Consistency: A consistent estimator is one which produces a better and better estimate as the sample size of the data increases.
- Bias: The bias of an estimator is the difference between estimator's expected value and the true value of the parameter being estimated.

15.3 Ordinary Least Squares

- A method of parameter estimation to determine the line of best fit to the data, where best fit is defined as the minimization sum of squared residuals.
- Least Squares Estimator

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

- Assumptions
 1. The regression model is linear in parameters.
 2. The samples are randomly sampled.
 3. The errors and independent variables are independent (errors are orthogonal to regressors).
 4. There is no multi-collinearity (there is no linear relationship between independent variables).

5. The variance of the errors of the regression model is constant (the errors are homoskedastic). $D = \sigma^2 \mathbf{I}$
6. The errors are not correlated with each other or themselves.
7. The errors are normally distributed.

15.4 Generalized Least Squares

- A generalization of the OLS estimation technique.
- It is suitable for fitting linear models on data that exhibit heteroskedasticity (non-constant variance) and/or autocorrelation, which are common properties in real world data. OLS assumes that these properties are not present, therefore making it unfit for estimation in this case.
- When errors are i.i.d. ($\Sigma = \sigma^2 \mathbf{I}$), OLS is appropriate, but when errors are not i.i.d. ($\Sigma \neq \sigma^2 \mathbf{I}$), OLS is not appropriate.
- Generalized Least Squares Estimator

$$\hat{\beta} = (\mathbf{X}^T \Sigma^{-1} \mathbf{X})^{-1} \mathbf{X}^T \Sigma^{-1} \mathbf{Y}$$

- A special case of GLS is Weighted Least Squares, which occurs when Σ is a diagonal matrix. The intuition is observations are weighted according to their volatilities and are estimated in the same manner as regular GLS, that is

$$\hat{\beta} = (\mathbf{X}^T \Sigma^{-1} \mathbf{X})^{-1} \mathbf{X}^T \Sigma^{-1} \mathbf{Y}$$

- GLS is theoretically more efficient than OLS, but hard to implement in practice.
- Drawbacks
 1. Σ is usually unknown and has to be estimated.

15.5 Heteroskedasticity and Autocorrelation Consistent (HAC) Standard Errors

- These are methods to estimate consistent standard errors when errors exhibit heteroskedasticity and/or autocorrelation.
- White Heteroskedasticity Consistent (HC) Covariance Estimator

$$\hat{\text{Cov}}_{HC}(\hat{\beta} | \mathbf{x}_1, \dots, \mathbf{x}_n) = (\mathbf{X}^T \mathbf{X})^{-1} \hat{\mathbf{C}}_{HC} (\mathbf{X}^T \mathbf{X})^{-1}$$

where

$$\hat{\mathbf{C}}_{HC} = \mathbf{X}^T \hat{\Sigma}_{\epsilon} \mathbf{X}$$

- Newey West Heteroskedasticity and Autocorrelation Consistent (HAC)

15.6 Maximum Likelihood

- **Likelihood Function:** Let X_1, \dots, X_n be an i.i.d. random sample from a distribution with a parameter θ (θ may be a real value or a vector, $\theta = (\theta_1, \dots, \theta_k)$). Suppose x_1, \dots, x_n are observed values of X_1, \dots, X_n . If the X_i 's are discrete r.v.s, the likelihood function is defined as the probability of observing the sample in the real world as a function of θ

$$L(x_1, \dots, x_n | \theta) = P(X_1 = x_1, \dots, X_n = x_n | \theta) = \prod_{x=1}^n P(x_i | \theta)$$

If the X_i 's are continuous r.v.s, we have a similar likelihood function

$$L(x_1, \dots, x_n | \theta) = f(X_1 = x_1, \dots, X_n = x_n | \theta) = \prod_{x=1}^n f(x_i | \theta)$$

- As sums are often easier to deal with than products, we modify this to the log likelihood function. If the X_i 's are discrete r.v.s, this is defined as

$$\log L(x_1, \dots, x_n | \theta) = \sum_{x=1}^n \log P(x_i | \theta)$$

If the X_i 's are continuous r.v.s, we have a similar log likelihood function

$$\log L(x_1, \dots, x_n | \theta) = \sum_{x=1}^n \log f(x_i | \theta)$$

- Naturally, this is a value we would like to maximize, leading us to maximum log likelihood estimation.

$$\hat{\theta}_{mle} = \max_{\theta} \log L(x | \theta)$$

15.7 MLE with Autocorrelation

- Since the sample is no longer i.i.d, the density is no longer the product of the individual densities. Instead, we use conditional densities

$$f(x_1, \dots, x_n) = f(x_n | x_{n-1}, \dots, x_1) f(x_{n-1} | x_{n-2}, \dots, x_1) \dots f(x_2 | x_1) f(x_1)$$

16 Multivariate Time Series

- Multivariate regression is not multiple regression.
- Multivariate regression refers to modeling multiple target variables at the same time.

16.1 Vector Autoregressions (VAR)

- A generalization of AR processes for multivariate series.
- Suppose that for each t

$$\mathbf{Y}_t = \begin{bmatrix} Y_{1,t} \\ Y_{2,t} \\ \vdots \\ Y_{d,t} \end{bmatrix}$$

is a d -dimensional multivariate time series.

- The definition of stationarity for multivariate time series is the same with univariate time series.
- A multivariate time series is said to be stationary if, for every n and m , $\mathbf{Y}_1, \dots, \mathbf{Y}_n$ and $\mathbf{Y}_{1+m}, \dots, \mathbf{Y}_{n+m}$ have the same distributions.

16.1.1 VAR(1)

- A bivariate VAR(1) process may be written

$$\begin{bmatrix} Y_{1,t} \\ Y_{2,t} \end{bmatrix} = \begin{bmatrix} \phi_{0,1} \\ \phi_{0,2} \end{bmatrix} + \begin{bmatrix} \phi_{11,1} & \phi_{12,1} \\ \phi_{21,1} & \phi_{22,1} \end{bmatrix} \begin{bmatrix} Y_{1,t-1} \\ Y_{2,t-1} \end{bmatrix} + \begin{bmatrix} \epsilon_{1,t} \\ \epsilon_{2,t} \end{bmatrix}$$

where

$$\Phi_0 = \begin{bmatrix} \phi_{0,1} \\ \phi_{0,2} \end{bmatrix}$$

and

$$\Phi_1 = \begin{bmatrix} \phi_{11,1} & \phi_{12,1} \\ \phi_{21,1} & \phi_{22,1} \end{bmatrix}$$

expanding this out, we have the following equations

$$Y_{1,t} = \phi_{0,1} + \phi_{11,1}Y_{1,t-1} + \phi_{12,1}Y_{2,t-1} + \epsilon_{1,t}$$

and

$$Y_{2,t} = \phi_{0,2} + \phi_{21,1}Y_{1,t-1} + \phi_{22,1}Y_{2,t-1} + \epsilon_{2,t}$$

where $\phi_{ij,1}$ is the amount of influence of $Y_{j,t-1}$ on $Y_{i,t}$.

- For a VAR(1) for k variables we have

$$\mathbf{Y}_t = \Phi_0 + \Phi_1 \mathbf{Y}_{t-1} + \epsilon_t$$

16.1.2 VAR(p)

- A VAR(p) may be written

$$\mathbf{Y}_t = \Phi_0 + \Phi_1 \mathbf{Y}_{t-1} + \dots + \Phi_p \mathbf{Y}_{t-p} + \epsilon_t$$

16.1.3 Cross-Correlation Function

- Suppose Y_j and Y_i are 2 component series of a stationary multivariate time series. The cross-correlation function (CCF) between Y_j and Y_i is defined as

$$\rho_{Y_i, Y_j}(h) = \text{Corr}\{Y_j(t), Y_i(t-h)\}$$

and is the correlation between Y_j at time t and Y_i at h time units earlier.

- Unlike the ACF, the CCF is not symmetric in the lag variable

$$\rho_{Y_j, Y_i}(h) \neq \rho_{Y_j, Y_i}(-h)$$

however what does hold true is the following

$$\rho_{Y_j, Y_i}(h) = \rho_{Y_i, Y_j}(-h)$$

16.1.4 Forecasting

- If mean is 0

$$E_t(\mathbf{Y}_{t+k}|\mathbf{Y}_t) = \Phi^k \mathbf{Y}_t$$

- If mean is non-zero

$$E_t(\mathbf{Y}_{t+k}|\mathbf{Y}_t) = \mu + \Phi^k(\mathbf{Y}_t - \mu)$$

16.2 Granger Causality

- A time series w_t Granger causes another time series y_t if w_t helps to forecast y_t , given past y_t .
- The idea is that w_t Granger causes y_t if w_t can assist in forecasting y_t by looking at past values of y_t .

17 Cointegration

17.1 Cointegration with Two Time Series

- Up until now, if we've had non-stationary time series, it has been a bad idea to regress them on one another.
- Sometimes we can find 2 or more non-stationary time series that are so closely connected that a linear combination of them is stationary. That is, they have a common trend, the linear combination gets rid of the trend, and the result is a stationary time series which is just the deviation from the trend.
- In these cases, we can create a cointegrating vector, which is the vector of the coefficients of this linear combination.
- Slightly more formally, 2 times series, $\mathbf{Y}_t = (Y_{1,t}, Y_{2,t})'$ are cointegrated if each is $I(1)$ but there exists a cointegration vector, $\alpha = (1, -\alpha)$, such that $\alpha' \mathbf{Y}_t = Y_{1,t} - \alpha Y_{2,t}$ is stationary. That is, $\alpha' \mathbf{Y}_t \sim I(0)$.

- In discussing cointegration, the key idea is that a series could be nonstationary, but at the same time, there could be some relationship in these variables that drive two or more of these variables at the same time.
- To test for cointegration, regress $Y_{1,t}$ on $Y_{2,t}$ to get an α and then we can check the residuals to see if they are stationary.
- The stationary cointegration error is the deviation from the trend as is defined

$$Z_t = Y_{1,t} - \alpha Y_{2,t}$$

18 Vector Error Correction Models

- The regression approach to cointegration is somewhat unsatisfactory, as one series must be chosen as the dependent variable, and this choice must be somewhat arbitrary.
- VECMs offer an alternate approach in which the deviation from the mean is called the error and whenever the stationary linear combination deviates from its mean, it is subsequently pushed back toward its mean (the error is corrected) to follow its long term path and maintain stationarity.
- Granger's Representation Theorem: The cointegrated series \mathbf{X}_t with cointegrating vector α has 2 equivalent representations:

1. Restricted VAR(p) in \mathbf{X}_t

$$\mathbf{X}_t = \Phi(L)\mathbf{X}_{t-1} + \mathbf{w}_t$$

where the coefficients of $\Phi(L)$ are subject to nonlinear restrictions.

2. VECM($p - 1$)

$$\Delta \mathbf{X}_t = \gamma z_{t-1} + \Psi(L)\Delta \mathbf{X}_{t-1} + \mathbf{w}_t$$

where $z_t = \alpha' \mathbf{X}_t$ and γ has at least one non-zero element. z_{t-1} specifies how much the variables will come back together and γ dictates how fast they come back together.

19 Present Value Relationships

- The expectation of the price at time t is the sum of all future dividends, discounted by the product of all interest rates in all future periods

$$P_t = E_t \left[\sum_{i=1}^{\infty} \frac{D_{t+i}}{\prod_{j=1}^i (1 + R_{t+j})} \right]$$

- Under Random Walk Hypothesis returns are not predictable.

19.1 Gordon Growth Model

-

$$D_t = \frac{1 + G}{R - G} D_t$$

- Assumptions

1. Dividends grow at rate G

$$D_{t+1} = (1 + G)D_t + e_{1,t+1}$$

2. Discount rate is R (constant return). This implies that, when defining P_t , $\prod_{j=1}^i (1 + R_{t+j}) \rightarrow (1 + R)^i$

- Price dividend ratio

$$\frac{P_t}{D_t} = \frac{1 + G}{R - G}$$

20 CAPM

21 Misc.

- Returns are closer to i.i.d. than prices and overall exhibit more attractive statistical qualities than prices, therefore making it more sensible to study returns rather than prices.

22 Definitions

- Homoskedasticity: A condition in which the variance of the residual is constant, that is, the error term does not vary much. In other words, the variance of the data points is roughly the same for all data points.